

Gaussian Processes: The Basics

CS698X: Topics in Probabilistic Modeling and Inference

Piyush Rai

Linear Models and Their Limitations

- Consider learning to map an input \mathbf{x} to the output y
- We've seen various discriminative models (linear and generalized linear models)


$$p(y|\mathbf{w}, \mathbf{x}) = \mathcal{N}(y|\mathbf{w}^\top \mathbf{x}, \beta^{-1}) \quad (\text{Linear Regression})$$

$$p(y|\mathbf{w}, \mathbf{x}) = [\sigma(\mathbf{w}^\top \mathbf{x})]^y [1 - \sigma(\mathbf{w}^\top \mathbf{x})]^{1-y} \quad (\text{Logistic Regression})$$


$$p(y|\mathbf{w}, \mathbf{x}) = \text{ExpFam}(\mathbf{w}^\top \mathbf{x}) \quad (\text{Generalized Linear Model})$$

Natural param of canonical GLM

- These have limited expressive power – can't learn nonlinear patterns



Nonlinear Regression



Nonlinear Classification



Learning Nonlinear Functions

- Assume the input to output relationship to be modeled by a nonlinear function f

$$p(y|f, \mathbf{x}) = \mathcal{N}(y|f(\mathbf{x}), \beta^{-1})$$

$$p(y|f, \mathbf{x}) = [\sigma(f(\mathbf{x}))]^y [1 - \sigma(f(\mathbf{x}))]^{1-y}$$

$$p(y|f, \mathbf{x}) = \text{ExpFam}(f(\mathbf{x}))$$

In all of these, the linear score $\mathbf{w}^T \mathbf{x}$ has been replaced by a nonlinear function $f(\mathbf{x})$



- Would like to model this function in a probabilistic/Bayesian manner
 - Nonlinearity + all the benefits of probabilistic/Bayesian modeling

- Some ways to achieve this

Example: Assuming \mathbf{x} is scalar,
 $\phi(\mathbf{x}) = [1, \mathbf{x}, \mathbf{x}^2, \dots, \mathbf{x}^k]$, for some k

- Ad-hoc: Manually define nonlinear features $\phi(\mathbf{x})$ + train Bayesian linear model
- Ad-hoc: Use a pre-train deep neural net to extract features $\phi(\mathbf{x})$ + train Bayesian linear model
- Bayesian Neural Networks (later)
- Gaussian Processes (a Bayesian approach to [kernel based](#) nonlinear learning; today)



Gaussian Process

Any choice of the GP covariance function has an associated feature map $\phi(x)$ for the inputs x

Hmmm.. So GPs look like kernel methods with all the benefits of probabilistic/Bayesian modeling



4

- A Gaussian Process (GP) defines a **distribution over functions** and is denoted as

Akin to how we define a Gaussian distribution over scalars/vectors, defined by a mean and variance/covariance matrix


Mean Function

Covariance Function

$$\mathcal{GP}(\mu(\cdot), \kappa(\cdot, \cdot))$$

Can also think of a function as an **infinite dimensional vector** of function's values at different inputs (x), i.e.,
 $f = [f(x_1), f(x_2), f(x_3), \dots]$

- Every draw/sample from $\mathcal{GP}(\mu, \kappa)$ will give a random function f



Each of these curves is a random function drawn from the GP

Mean Function $\mu(\cdot)$ defines the "average" function looks like:
 $\mu(x) = \mathbb{E}[f(x)]$

μ and κ can be pre-defined or can even be learned

Covariance Function $\kappa(\cdot, \cdot)$ defines **similarity between pairs of inputs** and controls the shape of these curves (also needed to be pos-sem-def)

- **IMP:** If $f \sim \mathcal{GP}(\mu, \kappa)$ then f 's value at any finite set of inputs is jointly Gaussian

Can concisely write it as

$$p(\mathbf{f}) = \mathcal{N}(\boldsymbol{\mu}, \mathbf{K})$$

$$p\left(\begin{bmatrix} f(x_1) \\ f(x_2) \\ \vdots \\ f(x_N) \end{bmatrix}\right) = \mathcal{N}\left(\begin{bmatrix} \mu(x_1) \\ \mu(x_2) \\ \vdots \\ \mu(x_N) \end{bmatrix}, \begin{bmatrix} \kappa(x_1, x_1) & \dots & \kappa(x_1, x_N) \\ \kappa(x_2, x_1) & \dots & \kappa(x_2, x_N) \\ \vdots & \ddots & \vdots \\ \kappa(x_N, x_1) & \dots & \kappa(x_N, x_N) \end{bmatrix}\right)$$

$N \times 1$ vector of f 's values: \mathbf{f}

$N \times 1$ mean vector: $\boldsymbol{\mu}$

$N \times N$ cov/kernel matrix (PSD): \mathbf{K}

Very useful property for **making predictions**: Knowing f 's value at some N "training" inputs, say, x_1, x_2, \dots, x_N , we can easily compute its value **at a new test input** x_* , using the Gaussian joint-to-conditional formula



Weight Space View vs Function Space View

- GPs are defined w.r.t. a **function space** that models input-output relationship
- In contrast, we have seen models that are defined w.r.t. a **weight space**, e.g.,

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \mathcal{N}(\mathbf{y}|\mathbf{X}\mathbf{w}, \beta^{-1}\mathbf{I}_N)$$

Likelihood

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0)$$

Prior over weight vector

$$p(\mathbf{y}|\mathbf{X}) = \int p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})d\mathbf{w} = \mathcal{N}(\mathbf{y}|\mathbf{X}\boldsymbol{\mu}_0, \beta^{-1}\mathbf{I}_N + \mathbf{X}\boldsymbol{\Sigma}_0\mathbf{X}^\top)$$

Marginal likelihood
after integrating
out the weights

$$p(\mathbf{y}|\mathbf{X}) = \mathcal{N}(\mathbf{y}|\mathbf{0}, \beta^{-1}\mathbf{I}_N + \mathbf{X}\mathbf{X}^\top)$$

Marginal likelihood assuming $\boldsymbol{\mu}_0 = \mathbf{0}$ and $\boldsymbol{\Sigma}_0 = \mathbf{I}$

$$p(\mathbf{y}|\mathbf{X}) = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{X}\mathbf{X}^\top)$$

Assuming noise-free likelihood

- Thus the joint distribution of the N responses $\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_N$ is a multivariate Gaussian

This equivalence also shows
that Bayesian linear regression
is a special case of GP with
linear kernel

$$p\left(\begin{bmatrix} \mathbf{y}_1 \\ \mathbf{y}_2 \\ \vdots \\ \mathbf{y}_N \end{bmatrix}\right) = \mathcal{N}\left(\begin{bmatrix} \mathbf{0} \\ \mathbf{0} \\ \vdots \\ \mathbf{0} \end{bmatrix}, \begin{bmatrix} \mathbf{x}_1^\top \mathbf{x}_1 & \dots & \mathbf{x}_1^\top \mathbf{x}_N \\ \mathbf{x}_2^\top \mathbf{x}_1 & \dots & \mathbf{x}_2^\top \mathbf{x}_N \\ \vdots & \ddots & \vdots \\ \mathbf{x}_N^\top \mathbf{x}_1 & \dots & \mathbf{x}_N^\top \mathbf{x}_N \end{bmatrix}\right)$$

Same as a GP $f(\mathbf{x}_i) = \mathbf{y}_i$, $\boldsymbol{\mu}(\mathbf{x}) = \mathbf{0}$ and linear
covariance/kernel function $\kappa(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^\top \mathbf{x}_j$

- Thus GPs can be seen as bypassing the weight space and directly defining the model using a marginal likelihood via a function space defined by the GP



Predicting using GP

The results we saw here relating the score f_* to \mathbf{f} will still hold ☺

We just need to use a likelihood model for y_n to handle such “noisy settings” (will see soon)

- We have already seen that

For example
 $p(y_n|f_n) = \mathcal{N}(y_n|f_n, \beta^{-1})$
 $p(y_n|f_n) = \text{Bernoulli}(y_n|\sigma(f_n))$

The setting considered on this slide is the “noiseless” setting where the response y_n is simply given by $y_n = f_n = f(x_n)$. More realistic settings will have each output y_n as a transformation of a “score” given by GP: $f_n = f(x_n)$

$$p\left(\begin{bmatrix} f(x_1) \\ f(x_2) \\ \vdots \\ f(x_N) \end{bmatrix}\right) = \mathcal{N}\left(\begin{bmatrix} \mu(x_1) \\ \mu(x_2) \\ \vdots \\ \mu(x_N) \end{bmatrix}, \begin{bmatrix} \kappa(x_1, x_1) & \dots & \kappa(x_1, x_N) \\ \kappa(x_2, x_1) & \dots & \kappa(x_2, x_N) \\ \vdots & \ddots & \vdots \\ \kappa(x_N, x_1) & \dots & \kappa(x_N, x_N) \end{bmatrix}\right) \xrightarrow{\text{concisely}} p(\mathbf{f}) = \mathcal{N}(\boldsymbol{\mu}, \mathbf{K})$$

- Let's assume the mean function $\mu(x) = 0$, thus $\boldsymbol{\mu} = \mathbf{0}$ and $p(\mathbf{f}) = \mathcal{N}(\mathbf{0}, \mathbf{K})$
- Assume we know $\mathbf{f} = [f(x_1), f(x_2), \dots, f(x_N)]$ and want to compute $f(x_*)$
- Due to the GP property, joint distribution of f 's values will always be Gaussian

$$p\left(\begin{bmatrix} \mathbf{f} \\ f_* \end{bmatrix}\right) = \mathcal{N}\left(\begin{bmatrix} \mathbf{0} \\ 0 \end{bmatrix}, \begin{bmatrix} \mathbf{K} & \mathbf{k}_* \\ \mathbf{k}_*^\top & \kappa(x_*, x_*) \end{bmatrix}\right)$$

where $\mathbf{k}_* = [\kappa(x_1, x_*), \kappa(x_2, x_*), \dots, \kappa(x_N, x_*)]^\top$

Important result

$(N+1) \times 1$ vector

$(N+1) \times (N+1)$ matrix

$N \times 1$ vector of similarities of x_* with each of the N training inputs

$$p(f_*|\mathbf{f}) = \mathcal{N}(\mathbf{k}_*^\top \mathbf{K}^{-1} \mathbf{f}, \kappa(x_*, x_*) - \mathbf{k}_*^\top \mathbf{K}^{-1} \mathbf{k}_*) = \mathcal{N}(\mu_*, \sigma_*^2)$$

Form of prediction similar to kernel methods but also get variances σ_*^2

- Exercise: Show that predictive mean $\mu_* = \sum_{i=1}^N \beta_i f_i = \sum_{i=1}^N \alpha_i \kappa(x_i, x_*)$


GP: A Visualization

- Assumed zero mean function and a squared exponential kernel

$$k_{\text{SE}}(x, x') = \sigma^2 \exp\left(-\frac{(x-x')^2}{2\ell^2}\right)$$


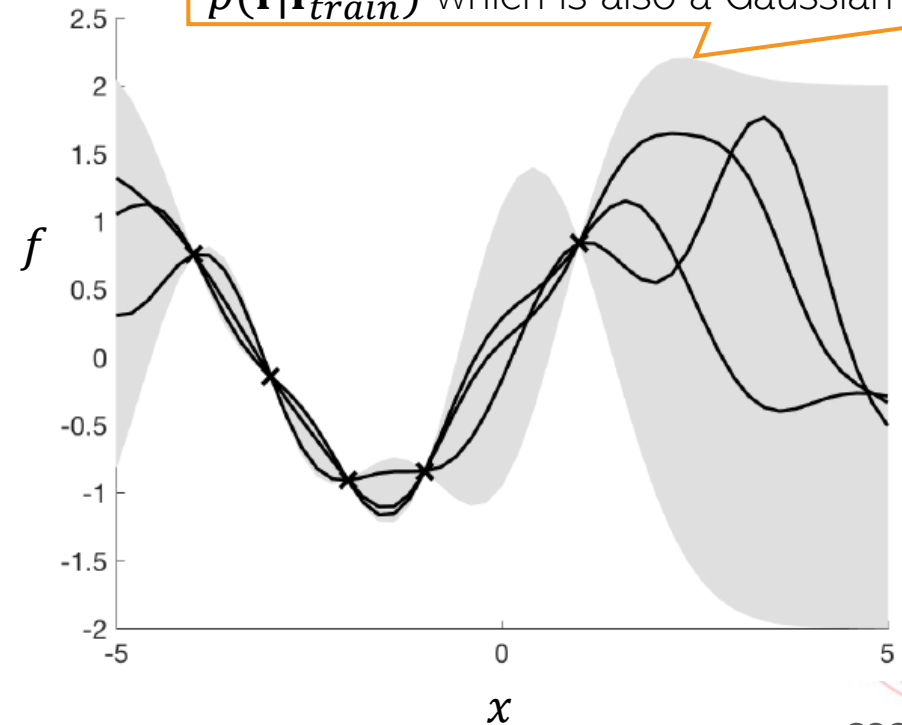
Each curve below is obtained by drawing a random \mathbf{f} from the GP "prior" $p(\mathbf{f}) = \mathcal{N}(\mathbf{0}, \mathbf{K})$ and plotting it.

\mathbf{K} is the kernel matrix of a finite number of inputs represented on the x axis (say 100 equi-spaced points between -5 and 5). \mathbf{f} will be a vector of f 's values at these inputs



Shaded area shows the predictive uncertainty for each of the test inputs (+/- 2 std)

Each curve below is obtained by drawing a random \mathbf{f} 's drawn from the GP posterior $p(\mathbf{f}|\mathbf{f}_{\text{train}})$ which is also a Gaussian



Coming Up

- GP for the “noisy” setting
 - Regression with Gaussian likelihood
 - Classification with Bernoulli/multinoulli likelihood
- Estimating the covariance/kernel function
- Connections with deep neural networks

$$p(y_n|f_n) = \mathcal{N}(y_n|f_n, \beta^{-1})$$
$$p(y_n|f_n) = \text{Bernoulli}(y_n|\sigma(f_n))$$

