

Gaussian Processes (Contd)

CS698X: Topics in Probabilistic Modeling and Inference

Piyush Rai

Recap: Gaussian Process

- A Gaussian Process (GP) defines a **distribution over functions** and is denoted as

$$\mathcal{GP}(\mu(.), \kappa(.,.))$$

- Assume training data with N input-output pairs

$$(x_1, f(x_1)), (x_2, f(x_2)), (x_3, f(x_3)), \dots (x_N, f(x_N))$$

- Assuming f has a GP prior $\mathcal{GP}(\mu(.), \kappa(.,.))$

Can concisely write it as

$$p(\mathbf{f}) = \mathcal{N}(\boldsymbol{\mu}, \mathbf{K})$$

$$p\left(\begin{bmatrix} f(x_1) \\ f(x_2) \\ \vdots \\ f(x_N) \end{bmatrix}\right) = \mathcal{N}\left(\begin{bmatrix} \mu(x_1) \\ \mu(x_2) \\ \vdots \\ \mu(x_N) \end{bmatrix}, \begin{bmatrix} \kappa(x_1, x_1) & \dots & \kappa(x_1, x_N) \\ \kappa(x_2, x_1) & \ddots & \kappa(x_2, x_N) \\ \vdots & \ddots & \vdots \\ \kappa(x_N, x_1) & \dots & \kappa(x_N, x_N) \end{bmatrix}\right)$$

- Assuming $\mu(.) = 0$, the prediction $f_* = f(x_*)$ for a test input x_* follows

$$p(f_* | \mathbf{f}) = \mathcal{N}(\mathbf{k}_*^\top \mathbf{K}^{-1} \mathbf{f}, \kappa(x_*, x_*) - \mathbf{k}_*^\top \mathbf{K}^{-1} \mathbf{k}_*) = \mathcal{N}(\mu_*, \sigma_*^2)$$

- The mean of this predictive distribution: $\mu_* = \sum_{i=1}^N \beta_i f_i = \sum_{i=1}^N \alpha_i \kappa(x_i, x_*)$
- The above setting is “noiseless”: Output is simply $f(x_n)$ with f modeled by a GP



GP for Noisy Setting

- Assume training data $(\mathbf{X}, \mathbf{y}) = \{(x_1, y_1), (x_2, y_2), (x_3, y_n), \dots (x_N, y_N)\}$
- Assume each y_n is obtained from $f_n = f(x_n)$ via an appropriate likelihood model, e.g.

$$\begin{aligned} p(y_n|f_n) &= \mathcal{N}(y_n|f_n, \beta^{-1}) \\ p(y_n|f_n) &= \text{Bernoulli}(y_n|\sigma(f_n)) \\ p(y_n|f_n) &= \text{ExpFam}(y_n|f_n) \end{aligned}$$

E.g., in the regression case, we can think of y_n as $y_n = f_n + \epsilon_n$ with $\epsilon_n \sim \mathcal{N}(0, \beta^{-1})$ being Gaussian noise

“Noise” simply refers to the fact that we are using a probability distribution to generate y_n and the probability distribution depends on a GP based “score” f_n



$N \times 1$ vector of the noisy outputs y_n 's in training data

$N \times 1$ vector of the “scores” f_n 's, each modeled by the GP

Note that in linear models, this score is simply $\mathbf{w}^T \mathbf{x}_n$

- Now we have a likelihood model $p(\mathbf{y}|\mathbf{f})$ for these “noisy” outputs
- Our “prior” over the function f is still a GP and is given by $p(\mathbf{f}) = \mathcal{N}(\mathbf{f}|\boldsymbol{\mu}, \mathbf{K})$
- IMP:** Prior $p(\mathbf{f})$ depends on training inputs \mathbf{X} (via \mathbf{K}) but not on outputs \mathbf{y}
- We can now combine the prior and likelihood to compute
 - GP posterior $p(\mathbf{f}|\mathbf{y})$, marginal likelihood $p(\mathbf{y})$, PPD $p(y_*|\mathbf{y})$, etc
 - Note: For Gaussian lik. based regression, PPD can be computed without computing $p(\mathbf{f}|\mathbf{y})$

It is almost similar to the noiseless case (will see shortly)

GP for Noisy Setting: Regression (Gaussian Lik.)

- Assume each response modeled by a Gaussian likelihood $p(y_n|f_n) = \mathcal{N}(y_n|f_n, \beta^{-1})$
- Denoting $\mathbf{y} = [y_1, y_2, \dots, y_N]$ and $\mathbf{f} = [f_1, f_2, \dots, f_N]$ and i.i.d. responses

$$p(\mathbf{y}|\mathbf{f}) = \mathcal{N}(\mathbf{y}|\mathbf{f}, \beta^{-1}\mathbf{I})$$

$p(\mathbf{f}|\mathbf{y})$ for general likelihoods
however will not be a Gaussian

Exercise: Derive by substituting the
prior and likelihood expressions

- Assume a zero-mean GP prior $p(\mathbf{f}) = \mathcal{N}(\mathbf{f}|\mathbf{0}, \mathbf{K})$
- For Gaussian likelihood, the posterior $p(\mathbf{f}|\mathbf{y}) \propto p(\mathbf{f}) p(\mathbf{y}|\mathbf{f})$ will be another Gaussian
- Posterior predictive $p(y_*|\mathbf{x}_*, \mathbf{y}, \mathbf{X})$ or $p(y_*|\mathbf{y})$ (skipping \mathbf{X}, \mathbf{x}_* from the notation)

Note: This PPD result is general
and holds for all likelihoods, not
just Gaussian

$$p(y_*|\mathbf{y}) = \int p(y_*|f_*)p(f_*|\mathbf{y})df_*$$

and known hyperparams of the
likelihood and GP prior

It's form will depend on
the likelihood model

$$p(f_*|\mathbf{y}) = \int p(f_*|\mathbf{f})p(\mathbf{f}|\mathbf{y})d\mathbf{f}$$

Always a
Gaussian for GP

GP posterior
(may or may
not be a
Gaussian)

- For Gaussian likelihood, PPD can be computed without using this general method
 - The form will be almost identical to the noiseless case (we will see shortly)



GP for Noisy Setting: Regression (Gaussian Lik.)

- For Gaussian lik, we can get PPD $p(\mathbf{y}_*|\mathbf{y})$ without computing the GP posterior $p(\mathbf{f}|\mathbf{y})$
- Note that, in this case, the marginal likelihood is also a Gaussian

A useful quantity when learning hyperparams of the GP covariance/kernel

$$p(\mathbf{y}) = \int p(\mathbf{y}|\mathbf{f})p(\mathbf{f})d\mathbf{f} = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{K} + \beta^{-1}\mathbf{I}_N) = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{C}_N)$$

- The joint distribution of the training \mathbf{y} and test response y_* is also a Gaussian

Note: All hyperparams assumed to be known

$$p\left(\begin{bmatrix} \mathbf{y} \\ y_* \end{bmatrix}\right) = \mathcal{N}\left(\begin{bmatrix} \mathbf{y} \\ y_* \end{bmatrix} \middle| \begin{bmatrix} \mathbf{0} \\ 0 \end{bmatrix}, \begin{bmatrix} \mathbf{C}_N & \mathbf{k}_* \\ \mathbf{k}_*^\top & \kappa(x_*, x_*) + \beta^{-1} \end{bmatrix}\right)$$

Identical to the noiseless case except the additional β^{-1} term on the diagonal

- Using the above, we can easily obtain $p(y_*|\mathbf{y})$ using Gaussian properties

$$p(y_*|\mathbf{y}) = \mathcal{N}(\mathbf{k}_*^\top \mathbf{C}_N^{-1} \mathbf{y}, \kappa(x_*, x_*) - \mathbf{k}_*^\top \mathbf{C}_N^{-1} \mathbf{k}_* + \beta^{-1}) = \mathcal{N}(\mu_*, \sigma_*^2)$$

Weighted average of the training responses


μ_* has a similar interpretation as in the noiseless case

- This is almost identical to the expression of $p(f_*|\mathbf{f})$ from the noiseless case except \mathbf{K} there is replaced by \mathbf{C}_N and extra β^{-1} term in variance



GP Regression: An Illustration

- The figure below shows GP predictive mean and variance as noise variance changes



Blue curve: True function
Red point: Training inputs (noisy)
Red curve: Learned predictive mean
Shaded region: ± 3 std-dev

- As expected, the predictive mean worsens and predictive variance increases as the noise variance increases



GP for Noisy Setting: Classification and GLM

- Binary classification: Now likelihood will be Bernoulli: $p(y_n|f_n) = \text{Bernoulli}(y_n|\sigma(f_n))$
- For multi-class ($K > 2$) GP, $p(y_n|f_n)$ will be multinoulli and f_n will be a $K \times 1$ vector
- For GP based GLM, $p(y_n|f_n)$ will be some exp-family distribution
- The prior $p(\mathbf{f})$ will still be a GP. Assuming a zero-mean GP prior $p(\mathbf{f}) = \mathcal{N}(\mathbf{f}|\mathbf{0}, \mathbf{K})$
- The posterior predictive $p(y_*|\mathbf{y})$ can again be written as

$$\begin{aligned} p(y_*|\mathbf{y}) &= \int p(y_*|f_*)p(f_*|\mathbf{y})df_* \\ &= \int p(y_*|f_*)p(f_*|\mathbf{f})p(\mathbf{f}|\mathbf{y})d\mathbf{f}df_* \end{aligned}$$

- This in general is not as easy to compute unlike the case of GP regression we saw
 - $p(f_*|\mathbf{f})$ is still not a problem (will be Gaussian due to the GP property)
 - GP posterior $p(\mathbf{f}|\mathbf{y}) \propto p(\mathbf{f})p(\mathbf{y}|\mathbf{f})$ will require approximation (Laplace, MCMC, variational, etc)
 - The overall integral will require approximation as well



Learning Hyperparameters in GP based Models

- Can learn the hyperparameters of the GP prior as well as of the likelihood model
- Assuming $\mu = \mathbf{0}$, the hyperparams of GP are cov/kernel function hyperparams

$$\kappa(\mathbf{x}_n, \mathbf{x}_m) = \exp\left(-\frac{\|\mathbf{x}_n - \mathbf{x}_m\|^2}{\gamma}\right) \quad \text{(RBF kernel)}$$

$$\kappa(\mathbf{x}_n, \mathbf{x}_m) = \exp\left(-\sum_{d=1}^D \frac{(\mathbf{x}_{nd} - \mathbf{x}_{md})^2}{\gamma_d}\right) \quad \text{(ARD kernel)}$$

$$\kappa(\mathbf{x}_n, \mathbf{x}_m) = \kappa_{\theta_1}(\mathbf{x}_n, \mathbf{x}_m) + \kappa_{\theta_2}(\mathbf{x}_n, \mathbf{x}_m) + \dots + \kappa_{\theta_M}(\mathbf{x}_n, \mathbf{x}_m) \quad \text{(flexible composition of multiple kernels)}$$

Can help in feature selection (irrelevant features will tend to have very large γ_d)

Different RBF kernel bandwidth γ_d for each feature

Ability to learn kernel hyperparams (without cross-valid) is another very appealing property of GP



- MLE-II is a popular choice for learning these hyperparams (otherwise MCMC, VI, etc)
- Denoting the covariance/kernel matrix as \mathbf{K}_θ , for Gaussian likelihood case, the marg-lik

$$p(\mathbf{y}|\theta, \beta^{-1}) = \mathcal{N}(\mathbf{y}|\mathbf{0}, \mathbf{K}_\theta + \beta^{-1}\mathbf{I}_N)$$

- This can be maximized to learn θ and β
- For **non-Gaussian likelihoods**, the marg-lik itself will need to be approximated



Coming Up

- Some aspects of GPs
 - Scalability
 - Connections with neural nets
 - Some recent advances

